



CMPS 460 – Spring 2022

MACHINE LEARNING

Tamer Elsayed




Image hosted by: WittySparks.com | Image source: Pixabay.com

5.a

Practical Issues: Dealing with Features



**Sec 2.7,
Sec. 3.1,
Sec 5-5.4**

Roadmap ...


- Learning algorithm is only one of many steps in designing a ML application
- Practical strategies:
 - Improving features
 - Evaluation
 - Cross Validation
 - Statistical Significance
 - Debugging
- Fundamental ML concepts: estimation vs. approximation error



ML Application ...

Typical Design Process for a ML App.

1	real world goal	increase revenue
2	real world mechanism	better ad display
3	learning problem	classify click-through
4	data collection	interaction w/ current system
5	collected data	query, ad, click
6	data representation	bow ² , \pm click
7	select model family	decision trees, depth 20
8	select training data	subset from april'16
9	train model & hyperparams	final decision tree
10	predict on test data	subset from may'16
11	evaluate error	zero/one loss for \pm click
12	deploy!	(hope we achieve our goal)




Dealing with Features

Features, values, vectors ...

- To a machine, the **features** themselves have no meaning.
 - Only the **feature values**!
- **Feature vector** $x = [x_1, x_2, \dots, x_D]$: consisting of one “dimension” for each feature, where each dimension is simply some real value.

[0.1 -3.4 5.0 1.0 0.0 4.1 0.69]



value of feature #3

Feature Types

*mapping in a
feature vector?*

- **Quantitative**

- Binary: 0/1, -1/1, T/F
- Real-valued: GPA, temperature, salary

- **Categorical**

- Sports, health, political, educational, economical
- Red, green, blue, yellow

- **Ordinal**

- Low, medium, high
- Weak, fair, strong, very strong
- Positive, neutral, negative

How about text?

- Raw text cannot be used directly in ML.

Why?

- Text is composed of a sequence of words.
- Words indicate meaning.

Words as features!

Text Representation: Bag of Words

- Simple representation of text used in NLP and IR.
- Text (e.g., sentence, tweet, article, web page) is represented as the **bag** (or a set) **of** its **words**, ignoring its order.
- ***Each word is a feature***
 - the feature value can be binary (appears in the text or not) or frequency of the word in the text (or even better – later).

Example

Given the following collection of documents:

d_1 : he likes to wink he likes to drink

d_2 : he likes to drink and drink and drink

d_3 : the thing he likes to drink is ink

d_4 : the ink he likes to drink is pink

d_5 : he likes to wink and drink pink ink

Feature vectors: binary values

	d_1	d_2	d_3	d_4	d_5
<i>he</i>	1	1	1	1	1
<i>likes</i>	1	1	1	1	1
<i>to</i>	1	1	1	1	1
<i>wink</i>	1	0	0	0	1
<i>drink</i>	1	1	1	1	1
<i>and</i>	0	1	0	0	1
<i>the</i>	0	0	1	1	0
<i>thing</i>	0	0	1	0	0
<i>ink</i>	0	0	1	1	1
<i>is</i>	0	0	1	1	0
<i>pink</i>	0	0	0	1	1

Vocabulary

Feature vectors: frequency values ...

	d_1	d_2	d_3	d_4	d_5
<i>he</i>	2	1	1	1	1
<i>likes</i>	2	1	1	1	1
<i>to</i>	2	1	1	1	1
<i>wink</i>	1	0	0	0	1
<i>drink</i>	1	3	1	1	1
<i>and</i>	0	2	0	0	1
<i>the</i>	0	0	1	1	0
<i>thing</i>	0	0	1	0	0
<i>ink</i>	0	0	1	1	1
<i>is</i>	0	0	1	1	0
<i>pink</i>	0	0	0	1	1

Vocabulary

Importance of Good Features ...

- “Garbage in, Garbage out”!
- Learning algorithms can’t compensate for useless training examples
 - e.g., if all features are irrelevant
- Feature design can have bigger impact on performance than tweaking the learning algorithm
 - e.g., feature combination

Irrelevant & Redundant Features

- **Irrelevant features:** completely uncorrelated with the prediction task.
 - e.g., the word “the”
- **Redundant features:** highly correlated, regardless of being relevant or not.
 - e.g., close pixels in images

DT?

kNN?

Perceptron?

Feature Pruning

- Very useful and applied in many applications.
- Easiest in the case of binary features.
 - If appears some small number K times
 - e.g., misspellings
 - If appears in all-but- K times
 - e.g., the word “the”
- For real-valued features
 - look for features with low variance.

Normalization

- Feature normalization

Centering:

$$x_{n,d} \leftarrow x_{n,d} - \mu_d$$

Variance Scaling:

$$x_{n,d} \leftarrow x_{n,d} / \sigma_d$$

Absolute Scaling:

$$x_{n,d} \leftarrow x_{n,d} / r_d$$

where:

$$\mu_d = \frac{1}{N} \sum_n x_{n,d}$$

$$\sigma_d = \sqrt{\frac{1}{N-1} \sum_n (x_{n,d} - \mu_d)^2}$$

$$r_d = \max_n |x_{n,d}|$$

Min-max Scaling


$$X_{norm} = \frac{X - X_{min}}{X_{max} - X_{min}}$$

- Example normalization

$$\mathbf{x}_n \leftarrow \mathbf{x}_n / ||\mathbf{x}_n||$$

Normalization: Centering

$$x_{n,d} \leftarrow x_{n,d} - \mu_d$$



<http://luisevalencia.com/feature-normalization/>


Normalization: Z-score normalization

Centering:

$$x_{n,d} \leftarrow x_{n,d} - \mu_d$$

Variance Scaling:

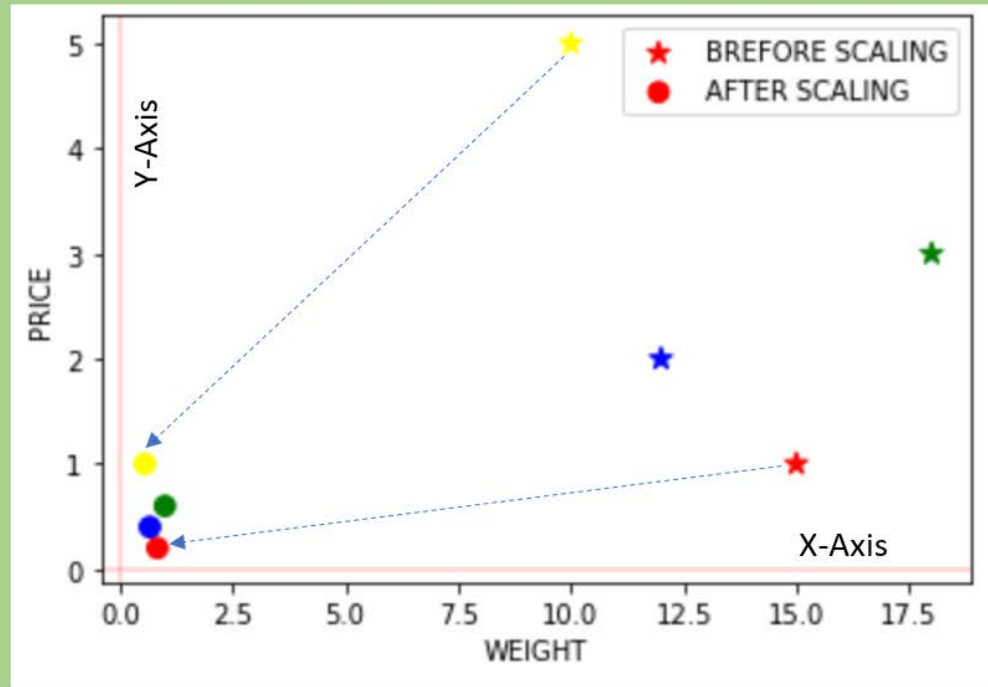
$$x_{n,d} \leftarrow x_{n,d} / \sigma_d$$



<https://towardsdatascience.com/all-about-feature-scaling-bcc0ad75cb35>

Normalization: Absolute Scaling


$$x_{n,d} \leftarrow x_{n,d} / r_d$$



<https://towardsdatascience.com/all-about-feature-scaling-bcc0ad75cb35>

Normalization: Min-max Scaling


$$X_{norm} = \frac{X - X_{min}}{X_{max} - X_{min}}$$



<https://towardsdatascience.com/all-about-feature-scaling-bcc0ad75cb35>

Normalization: Example Normalization

$$x_n \leftarrow x_n / ||x_n||$$



<https://towardsdatascience.com/all-about-feature-scaling-bcc0ad75cb35>

Logarithmic Transformation

- Mainly for textual features.
- Importance is not linear with frequency

$$x_d \mapsto \log_2(x_d + 1)$$

Feature Combination

- DT is essentially building meta features.

How?

- Constructing meta features for perceptron from DT
 1. train a shallow DT to extract meta features
 2. add only those feature combinations to the feature set for the perceptron.