



CMPS 460 – Spring 2022

MACHINE LEARNING

Tamer Elsayed




Image hosted by: WittySparks.com | Image source: Pixabay.com

7.c

Linear Models: Gradient Descent



Chapter 7:
7.4

Optimization Framework

**Objective
function**

Loss function
measures how well
classifier fits training
data

Regularizer
prefers solutions
that generalize
well

$$\min_{\mathbf{w}, b} L(\mathbf{w}, b) = \min_{\mathbf{w}, b} \sum_{n=1}^N \mathbb{I}(y_n(\mathbf{w}^T \mathbf{x}_n + b) < 0) + \lambda R(\mathbf{w}, b)$$

$l(y_n, \hat{y}_n)$

λ : parameter that controls
the importance of the
regularization term

- Different loss function approximations
 - easier to optimize
- Regularizer
 - prevents overfitting/prefers simple models.



Optimization with Gradient Descent


Gradient Descent

- A general solution for our optimization problem

$$\min_{\mathbf{w}, b} L(\mathbf{w}, b) = \min_{\mathbf{w}, b} \sum_{n=1}^N \mathbb{I}(y_n(\mathbf{w}^T \mathbf{x}_n + b) < 0) + \lambda R(\mathbf{w}, b)$$


$l(y_n, \hat{y}_n)$

Climbing a hill!



Gradient-based Optimization

- Suppose the goal is to find maximum of $f(\mathbf{x})$.
 1. The optimizer maintains a current estimate of the parameter of interest, \mathbf{x} .
 2. At each step, measure the **gradient \mathbf{g}** of $f(\mathbf{x})$ at the current location, \mathbf{x} .
 3. Then take a step in the direction of the gradient, where the size of the step is controlled by η .



$$\mathbf{x} \leftarrow \mathbf{x} + \eta \mathbf{g} \text{ (Gradient Ascent)}$$

- **Gradient Descent:** opposite of gradient ascent.

$$\mathbf{x} \leftarrow \mathbf{x} - \eta \mathbf{g} \text{ (Gradient Descent)}$$

Gradient Descent

- A general solution for our optimization problem

$$\min_{\mathbf{w}, b} L(\mathbf{w}, b) = \min_{\mathbf{w}, b} \sum_{n=1}^N \mathbb{I}(y_n(\mathbf{w}^T \mathbf{x}_n + b) < 0) + \lambda R(\mathbf{w}, b)$$

$l(y_n, \hat{y}_n)$

Idea: take iterative steps to update parameters
in the direction of the gradient

Gradient Descent in 1-Dim




Image source: <https://hackernoon.com/gradient-descent-aynk-7cbe95a778da>

Gradient Descent in 2-Dim




Image source: <https://towardsdatascience.com/gradient-descent-in-a-nutshell-eaf8c18212f0>

Gradient Descent Algorithm

Objective function
to minimize

Number of steps

Step size


Algorithm 22 GRADIENTDESCENT($\mathcal{F}, K, \eta_1, \dots$)

```


1:  $\mathbf{z}^{(0)} \leftarrow \langle 0, 0, \dots, 0 \rangle$  // initialize variable we are optimizing
2: for  $k = 1 \dots K$  do
3:    $\mathbf{g}^{(k)} \leftarrow \nabla_{\mathbf{z}} \mathcal{F}|_{\mathbf{z}^{(k-1)}}$  // compute gradient at current location
4:    $\mathbf{z}^{(k)} \leftarrow \mathbf{z}^{(k-1)} - \eta^{(k)} \mathbf{g}^{(k)}$  // take a step down the gradient
5: end for
6: return  $\mathbf{z}^{(K)}$ 

```

Impact of Step Size



Too small: converge
very slowly



Too big: overshoot and
even diverge

Image source: <https://towardsdatascience.com/gradient-descent-in-a-nutshell-eaf8c18212f0>

Gradient Descent

When to stop?


- When the gradient gets close to zero
- When the objective stops changing much
- When the parameters stop changing much
- When performance on held-out dev set plateaus

How to choose the step size?

- Constant
- Better: start with large steps, then take smaller steps





Example

$$f(x) = \frac{x^2 \cos(x) - x}{10}$$



Can we apply GD?

Example




Exercise

$$x^4 + 3x^3 - 2x$$

$$x_0 = 1, \alpha = 0.2$$

$$x_0 = 1, \alpha = 1.5$$



Gradients for Multivariate Objectives

- Consider the following learning objective

$$\mathcal{L}(\boldsymbol{w}, b) = \sum_n \exp \left[-y_n (\boldsymbol{w} \cdot \boldsymbol{x}_n + b) \right] + \frac{\lambda}{2} \|\boldsymbol{w}\|^2$$

What do we need to do to run gradient descent?

(1) Derivative with respect to b

$$\frac{\partial \mathcal{L}}{\partial b} = \frac{\partial}{\partial b} \sum_n \exp [-y_n(\mathbf{w} \cdot \mathbf{x}_n + b)] + \frac{\partial}{\partial b} \frac{\lambda}{2} \|\mathbf{w}\|^2 \quad (6.12)$$

$$= \sum_n \frac{\partial}{\partial b} \exp [-y_n(\mathbf{w} \cdot \mathbf{x}_n + b)] + 0 \quad (6.13)$$

$$= \sum_n \left(\frac{\partial}{\partial b} - y_n(\mathbf{w} \cdot \mathbf{x}_n + b) \right) \exp [-y_n(\mathbf{w} \cdot \mathbf{x}_n + b)] \quad (6.14)$$


$$= - \sum_n y_n \exp [-y_n(\mathbf{w} \cdot \mathbf{x}_n + b)] \quad (6.15)$$

(2) Gradient with respect to w

$$\nabla_w \mathcal{L} = \nabla_w \sum_n \exp [-y_n(\mathbf{w} \cdot \mathbf{x}_n + b)] + \nabla_w \frac{\lambda}{2} \|\mathbf{w}\|^2 \quad (6.16)$$

$$= \sum_n (\nabla_w - y_n(\mathbf{w} \cdot \mathbf{x}_n + b)) \exp [-y_n(\mathbf{w} \cdot \mathbf{x}_n + b)] + \lambda \mathbf{w} \quad (6.17)$$

$$= - \sum_n y_n \mathbf{x}_n \exp [-y_n(\mathbf{w} \cdot \mathbf{x}_n + b)] + \lambda \mathbf{w} \quad (6.18)$$



Subgradients

Subgradients


- Problem: some objective functions are not differentiable everywhere
 - e.g., Hinge loss, l_1 norm

Solution: Subgradient Optimization

- Let's ignore the problem, and just try to apply gradient descent anyway!!
- We will just differentiate by parts...

Subgradients

- Generalizes the derivative to functions which are not differentiable.
- For any x_0 in the domain of the function, one can draw a line which goes through the point $(x_0, f(x_0))$ and which is everywhere either touching or below the graph of f .
- Set-valued



Example: Subgradient of Hinge Loss

- For a given example n

$$\partial_w \max\{0, 1 - y_n(\mathbf{w} \cdot \mathbf{x}_n + b)\} \quad (7.24)$$

$$= \partial_w \begin{cases} 0 & \text{if } y_n(\mathbf{w} \cdot \mathbf{x}_n + b) > 1 \\ 1 - y_n(\mathbf{w} \cdot \mathbf{x}_n + b) & \text{otherwise} \end{cases} \quad (7.25)$$

$$= \begin{cases} \partial_w 0 & \text{if } y_n(\mathbf{w} \cdot \mathbf{x}_n + b) > 1 \\ \partial_w 1 - y_n(\mathbf{w} \cdot \mathbf{x}_n + b) & \text{otherwise} \end{cases} \quad (7.26)$$

$$= \begin{cases} 0 & \text{if } y_n(\mathbf{w} \cdot \mathbf{x}_n + b) > 1 \\ -y_n \mathbf{x}_n & \text{otherwise} \end{cases} \quad (7.27)$$

Subgradient Descent for Hinge Loss

Algorithm 23 HINGEREGULARIZEDGD($\mathbf{D}, \lambda, \text{MaxIter}$)

```

1:  $\mathbf{w} \leftarrow \langle 0, 0, \dots, 0 \rangle$  ,  $b \leftarrow 0$  // initialize weights and bias
2: for  $iter = 1 \dots \text{MaxIter}$  do
3:    $\mathbf{g} \leftarrow \langle 0, 0, \dots, 0 \rangle$  ,  $g \leftarrow 0$  // initialize gradient of weights and bias
4:   for all  $(x, y) \in \mathbf{D}$  do
5:     if  $y(\mathbf{w} \cdot \mathbf{x} + b) \leq 1$  then
6:        $\mathbf{g} \leftarrow \mathbf{g} + y \mathbf{x}$  // update weight gradient
7:        $g \leftarrow g + y$  // update bias derivative
8:     end if
9:   end for
10:   $\mathbf{g} \leftarrow \mathbf{g} - \lambda \mathbf{w}$  // add in regularization term
11:   $\mathbf{w} \leftarrow \mathbf{w} + \eta \mathbf{g}$  // update weights
12:   $b \leftarrow b + \eta g$  // update bias
13: end for
14: return  $\mathbf{w}, b$ 

```

Anything wrong here?

Perceptron ...

Algorithm 5 PERCEPTRONTRAIN(\mathbf{D} , $MaxIter$)

```

1:  $w_d \leftarrow 0$ , for all  $d = 1 \dots D$                                 // initialize weights
2:  $b \leftarrow 0$                                                     // initialize bias
3: for  $iter = 1 \dots MaxIter$  do
4:   for all  $(x, y) \in \mathbf{D}$  do
5:      $a \leftarrow \sum_{d=1}^D w_d x_d + b$                                 // compute activation for this example
6:     if  $ya \leq 0$  then
7:        $w_d \leftarrow w_d + yx_d$ , for all  $d = 1 \dots D$                 // update weights
8:        $b \leftarrow b + y$                                             // update bias
9:     end if
10:  end for
11: end for
12: return  $w_0, w_1, \dots, w_D, b$ 

```

What is it optimizing?

Summary: Gradient Descent

- A generic algorithm to minimize objective functions
- Works well as long as functions are well behaved (i.e., convex)
- Subgradient descent can be used at points where derivative is not defined.
- Can we do better?
 - For some objectives, we can find closed form solutions (see CIML 7.6)